

Neighborhoods, Ripley's K, etc.

Gabriel Arellano

2026-05-25

Contents

1	Introduction	1
2	List of old and new functions	2
3	Code	3
4	Example of code behaviour	12
5	Example of a practical application	17
6	The ring case	19
7	Comparing observations with expectations	19
7.1	Example of Ripley's K	19
7.2	Example of numerical null model	20

1 Introduction

This is updated code to calculate abundances of neighbors at different scales, in circles and rings. It also contains instructions to compare observations with expectations, including Ripley's K and Besag's L.

Spatial neighborhood analyses are used to quantify how organisms, stems, species, functional groups, or other ecological units are arranged around focal individuals or locations. In forest ecology, these analyses are commonly used to evaluate local crowding, conspecific and heterospecific aggregation, spatial segregation, density dependence, habitat association, biomass accumulation around focal trees, and departures from spatial null expectations. Classical

summaries include neighbor counts or densities within circles or rings, Ripley's K, and transformations such as Besag's L.

This tutorial introduces updated code and instructions for calculating neighborhood summaries at multiple spatial scales. The workflow separates several operations that were previously bundled together: identifying neighbors within a distance threshold, defining focal and reference classes, handling circles and annuli, estimating available neighborhood area near valid boundaries (of arbitrary shape, but typically the rectangular plot limits), summarizing abundance as counts of points or other metrics, and comparing observations with null expectations. This separation is intentional.

Relative to the previous CTFS functions, the updated code is designed to be less tied to specific assumptions about rectangular plots, point classes, or abundance metrics. An important design criterion is to leave outside the functions everything that may require high flexibility by the potential users:

- filters / subsetting routines
- aggregating or summarizing functions
- null expectations

Below, there are examples on how to implement these things in a script.

Edge effects can be important. Handling boundaries and calculating areas is *way* more difficult than it seems. Much of the code below was developed to handle irregular boundaries, thus make the code as general as possible. It does not require rectangular boundaries and it does not use analytical corrections for edge effects. It directly uses a dense grid that count “points inside” to estimate the areas of irregular shapes (neighborhoods near possibly irregular boundaries). The density of the grid is adaptive, depending on the error tolerance. The functions include useful defaults, so you should be fine if you have a regular plot and just pass {x, y} coordinates of all the trees.

The code takes responsibility for the boundary- and area-related difficult calculations and for the *very fast* identification of pairs of neighbors. Filtering and summarizing (e.g. counting neighbors, calculating crowding indexes, etc.) will be responsibility of the user. Below there are useful templates or instructions.

The approach taken is very fast but demands a lot of memory, so this can be a practical limitation for large datasets. For the moment, I will wait for users' feedback.

2 List of old and new functions

The code here substitutes these original CTFS functions:

- “CalcRingArea” in `spatial > RipUvK.r`

- “partialcirclearea” in `spatial > RipUvK.r`
- “circlearea” in `spatial > RipUvK.r`
- “Annuli” in `spatial > RipUvK.r`
- “NDcount” in `spatial > NeighborDensityFun.r`
- “NeighborDensities” in `spatial > NeighborDensityFun.r`
- “RipUvK” in `spatial > RipUvK.r`

The new functions are:

- `fast_distances_up_to_r`: the workhorse to calculate distances from one set of focal points (query) to their neighbors (taken from a reference dataset), within a given threshold. This takes advantage of the very fast `dbscan::frNN` function, and will be called often from within other functions here.
- `boundary_around_xy`: an auxiliary function to estimate a (complex) boundary around a given set of points. This may or may not help much, but it is needed as a default in case the user does not provide a proper boundary.
- `core_vs_edge`: an auxiliary function that classifies a set of points into “out” (outside the provided boundary of arbitrary shape), “edge” (inside the boundary, but closer than a given distance to it), and “core” (inside the boundary and far from it). The main goal of this auxiliary function is to save computational costs by avoiding numerical calculations for the area of full circles around “core” points.
- `make_grid_inside_boundary`: creates a regular grid within the provided boundary. The grid will be then used to count number of points within shapes, so it is a very general way of estimating areas of truncated circles near the (possibly complex) boundary. It might have independent use in some applications.
- `guess_sufficient_grid_resolution`: the grid approach is flexible but, as it is a numerical estimate, it also carries some error. The `guess_sufficient_grid_resolution` estimates the resolution of the grid required to keep error in circle area estimates lower than a given tolerance (1% by default).
- ...

Filters, abundance metrics, null expectations, and classes to compare, are all outside these functions.

3 Code

The main function here is a function to calculate, fast, distances up to certain threshold. This is useful for the distances themselves and to identify neighbors.

```
# Fast calculation of distances up to a given radius:
fast_distances_up_to_r <- function(query, reference = NULL, radius, approx = 0) {
```

```

# Parameters:
# query: the set of focal points, around which we are going to look
# reference: the set of potential neighbors; will default to be the same as "query"
# radius: the size of the local neighborhood around the focal point
# approx: precision parameter to be passed to dbscan::frNN

# Reshape input
if(is.null(reference)) reference <- query
query <- as.matrix(query)
reference <- as.matrix(reference)

storage.mode(query) <- "double"
storage.mode(reference) <- "double"

# Calculate distances
nn <- dbscan::frNN(
  x = reference,
  query = query,
  eps = radius,
  sort = FALSE,
  approx = approx
)

# Reorganize output
k <- lengths(nn$id)

data.table::data.table(
  query_i = rep(seq_along(k), k),
  reference_j = unlist(nn$id, use.names = FALSE),
  dist = unlist(nn$dist, use.names = FALSE)
)[dist <= radius]
}

```

Before other functions, let's talk about boundaries. Edge effects are important: the available area for neighborhoods is smaller near the boundary than far from it. Here we try to make the code as general as possible, and a major improvement would be to allow boundaries of arbitrary shape. The most common case, however, would be the rectangular boundary. In the updated functions, the user can define boundaries as:

- sfc polygon or multipolygon object, for boundaries of any shape. Typical parameter name: "boundary"
- as rectangular windows, expressed as a vector of (xmin, ymin, xmax, ymax) plot limits. Typical parameter name: "window"

If the user is not explicit, the functions will guess the boundary by creating a complex boundary around the set of $\{x, y\}$ points. It can be handy if one has trees classified into habitats and want to restrict the analyses to within particular habitats, for example. This is a proper default but it will take time to compute, so in general it will always be better to know and define the boundary beforehand, before entering costly loops, etc.

```
get_boundary <- function(boundary = NULL, window = NULL, xy = NULL, q = 5)
{
  # Minimal for spatial objects
  if (!is.null(boundary)) {
    if (inherits(boundary, "sfg")) {
      boundary <- sf::st_sfc(boundary)
    } else {
      boundary <- sf::st_geometry(boundary)
    }
  }

  # From "window" to proper sf object
  if (is.null(boundary) && !is.null(window)) {
    xmin <- window[1]
    ymin <- window[2]
    xmax <- window[3]
    ymax <- window[4]

    boundary <- sf::st_sfc(
      sf::st_polygon(list(rbind(
        c(xmin, ymin),
        c(xmax, ymin),
        c(xmax, ymax),
        c(xmin, ymax),
        c(xmin, ymin)
      )))
    )
  }

  # If just the points are provided, guess the boundary
  if (is.null(boundary) && is.null(window)) {
    xy <- as.matrix(xy)
    stopifnot(ncol(xy) == 2L, nrow(xy) >= 2L)
    storage.mode(xy) <- "double"

    knn <- FNN::get.knn(xy, k = 1)
  }
}
```

```

buffer <- q * median(knn$nn.dist[, 1], na.rm = TRUE)
delta <- buffer / 10

cell_x <- floor(xy[, 1] / delta)
cell_y <- floor(xy[, 2] / delta)
cell_id <- paste(cell_x, cell_y, sep = "_")
xy <- xy[!duplicated(cell_id), , drop = FALSE]

pts <- sf::st_as_sf(
  as.data.frame(xy),
  coords = 1:2
)

boundary <- sf::st_buffer(pts, dist = buffer + delta)
boundary <- sf::st_union(boundary)
boundary <- sf::st_simplify(boundary, dTolerance = delta, preserveTopology = TRUE)
}

if (is.null(boundary)) {
  stop("Provide one of: boundary, window, or xy.")
}

# Final polygon cleanup
boundary <- sf::st_geometry(boundary)
boundary <- sf::st_make_valid(boundary)
boundary <- sf::st_collection_extract(boundary, "POLYGON")

if (length(boundary) == 0 || all(sf::st_is_empty(boundary))) {
  stop("Boundary does not contain polygon geometry.")
}

boundary <- sf::st_union(boundary)
boundary <- sf::st_make_valid(boundary)
boundary
}

```

Now, a useful auxiliary function to split a given set of points into “out”, “edge” and “core”, given a boundary and a distance threshold. It takes advantage of the fast calculation of distances (function above). This function has, at least, these two applications in the workflow:

- for the numerical estimates of areas of irregular shapes, it focuses on the points and grid that are near the boundary, using the formula for the circle area for points far from the

boundary.

- for the *error estimate* in those numerical estimates of area, it focuses on the core points instead: it makes the same calculation twice, and compares the known area of circles with the area derived from the grid.

```
# Classify points based on a given boundary.
# The boundary can be provided as a sfc_POLYGON of arbitrary shape
# or a rectangular window (xmin, ymin, xmax, ymax). If it is not
# provided it will be assumed rectangular around the points provided.

core_vs_edge <- function(xy,
                        boundary = NULL,
                        window = NULL,
                        radius) {

  # Parameters:
  # xy: a matrix or data.frame with x and y coordinates in 1st and 2nd column respectively
  # boundary: boundary defined as a sfc_POLYGON object, for boundaries of any shape
  # window: alternative to "boundary" for rectangular boundaries, as c(xmin, ymin, xmax, ymax)
  # radius: the distance to the boundary that separates core from edge points
  # assume_inside: will accelerate if we know all xy points are inside the boundary

  # Confirm 2-columns matrix for the point coordinates
  xy <- as.matrix(xy)
  stopifnot(ncol(xy) == 2L)
  storage.mode(xy) <- "double"

  # Confirm or guess boundary:
  boundary <- get_boundary(boundary = boundary, window = window, xy = xy)

  # Determine which points are inside:
  pts <- sf::st_as_sf(
    as.data.frame(xy),
    coords = 1:2,
    crs = sf::st_crs(boundary),
    remove = FALSE
  )
  inside <- lengths(sf::st_intersects(pts, boundary)) > 0

  # Densify boundary once, then use fast_distances_up_to_r()
  boundary <- sf::st_boundary(boundary)
  boundary <- sf::st_segmentize(boundary, dfMaxLength = radius / 10)
```

```

boundary_xy <- sf::st_coordinates(boundary)[, c("X", "Y"), drop = FALSE]
boundary_xy <- unique(boundary_xy)
storage.mode(boundary_xy) <- "double"

# Identify points near the boundary:
near <- rep(FALSE, nrow(xy))
ii <- which(inside)
if (length(ii)) {
  d <- fast_distances_up_to_r(
    query = xy[ii, , drop = FALSE],
    reference = boundary_xy,
    radius = 1.05 * radius # small safety margin
  )

  near[ii[unique(d$query_i)]] <- TRUE
}

# Build output into three classes:
location <- rep(NA, length(inside))
location[!inside] <- "out"
location[inside & near] <- "edge"
location[inside & !near] <- "core"
location
}

```

The neighborhood analyses are scale-explicit, looking at circles or rings of increasing sizes. The user may not have a predefined set of scales at which to measure the metrics of interest. This function implements useful defaults, given a set of query points:

```

get_radii <- function(query, rmax = NA, m = 512, case = c("circles", "rings"))
{
  # Default case: circles
  if(length(case) != 1) case = "circles"

  # Implements same defaults as Ripley's K function in spatstat
  if(is.na(rmax))
  {
    a = diff(range(query[,1], na.rm = TRUE))
    b = diff(range(query[,2], na.rm = TRUE))
    lambda = nrow(query) / (a * b)
    rmax = min(c(a/4, b/4, sqrt(1000 / (pi * lambda))))
  }
}

```



```

# Default for circles: simple equal increments in radii
radii.c <- seq(from = rmax/m, to = rmax, by = rmax/m)

# Default for rings: rings of equal area
# (It returns the outer radii: the inner radii is 0 and the same sequence but shifted)
radii.r <- rmax*sqrt((1:m)/m)

# Output
if(case == "circles") radii <- radii.c
if(case == "rings")   radii <- radii.r
radii
}

```

The fundamental piece of information to control for edge effects are the areas of the neighborhoods at different scales, both for the “core” points (easy) and the “edge” points (difficult). The code can handle irregular shapes, estimating areas numerically within a tolerance to error, defined by the user. To do so, it creates dense grids (but not excessively dense) within the valid boundary, and then counts points in the grid within the provided radii. The function below integrates useful defaults and rules of thumb.

```

get_area_of_neighborhoods <- function(query = NULL,
                                       boundary = NULL,
                                       window = NULL,
                                       radii = NULL,
                                       rmax = NA,
                                       case = c("circles", "rings"),
                                       tol = 0.01)
{
  # query: The set of focal or query points.
  #       2-columns table with x and y coordinates in 1st and 2nd column, respectively.
  # boundary: boundary defined as a sfc_POLYGON object, for boundaries of any shape
  # window: rectangular, alternative to "boundary", as c(xmin, ymin, xmax, ymax)
  # radii: defines the scales of the circular neighborhoods.
  # rmax: maximum radius for neighborhoods, if radii not provided
  # m: number of scales to evaluate, if radii not provided
  # case: circular vs. ring-shaped neighborhoods: will affect area and default radii
  # tol: the accepted error in area, 1% by default.

  # Default case: circles
  if(length(case) != 1) case = "circles"

```

```

# Confirm or guess boundary
boundary <- get_boundary(boundary = boundary, window = window, xy = query)
bb <- sf::st_bbox(boundary)
xmin = bb["xmin"]
xmax = bb["xmax"]
ymin = bb["ymin"]
ymax = bb["ymax"]

# Define radii, if the user has not provided them yet:
if(is.null(radii))
  radii <- get_radii(query = query, rmax = rmax, m = m, case = case)

# The areas of neighborhoods at different scales will
# repeat similar calculations, and will use different
# grids for each scale. The alternative is to commit
# to a very dense grid with low error in the worst case
# scenario, but costly calculations when it is not really
# needed to achieve a given tolerance.

# This runs a loop across scales. Before the loop,
# it creates empty objects to store results.
areas <- matrix(NA, nrow = nrow(query), ncol = m) # point x scale matrix
steps <- rep(NA, m)
errors <- rep(NA, m)
for(i in 1:m)
{
  # The target scale
  radius = radii[i]
  A0 = pi*radius^2

  # Core vs. edge points in the query:
  ceo_query <- core_vs_edge(xy = query, boundary = boundary, radius = radius)

  # Initialize coarse grid, then refine until error < tolerance.
  # This uses the core points in the query.
  step = min(c((xmax - xmin)/3, (ymax - ymin)/3))
  error = 999
  while(error > tol)
  {
    step = step / 2 # smaller step
    grid <- expand.grid(x = seq(from = xmin - step/2, to = xmax + step/2, by = step),
                       y = seq(from = ymin - step/2, to = ymax + step/2, by = step))
  }
}

```

```

pts <- sf::st_as_sf(grid, coords = c("x", "y"), remove = FALSE)
inside <- lengths(sf::st_intersects(pts, boundary)) > 0
grid <- grid[inside,, drop = FALSE]

# Do the estimate using the core points only:
d.test <- fast_distances_up_to_r(query = query[ceo_query == "core",],
                                reference = grid,
                                radius = radius)

A.estimated <- pmin(A0,
                    tabulate(d.test$query_i, nbins = sum(ceo_query == "core")) * step^2)

error = mean(abs(as.numeric(A.estimated) - A0) / A0)
}

# The error just calculated applies to the core area, but,
# as a general rule, error will be larger near the boundary.
# So, to reach the same error near the boundary, we need
# a greater grid resolution. This is a conservative heuristic
# derived from Steiner / parallel-body argument:
K = 1 + 4 / pi
step = step / K

steps[i] <- step
errors[i] <- error

# For reference, other rules of thumb or analytical results.
#step = tol * radius / (2 * sqrt(2)) # to constrain maximum error
#step = radius * (2 * tol)^(2/3) # to constrain average error

# With the updated step, we address the calculation of areas.
# This focuses on the edge points in the query.
# Note the 2r distance needed to cover the entire neighborhood of query
# points that are at r distance from the edge.

grid <- expand.grid(x = seq(from = xmin - step/2, to = xmax + step/2, by = step),
                   y = seq(from = ymin - step/2, to = ymax + step/2, by = step))

pts <- sf::st_as_sf(grid, coords = c("x", "y"), remove = FALSE)
inside <- lengths(sf::st_intersects(pts, boundary)) > 0
grid <- grid[inside,, drop = FALSE]

```

```

ceo_grid <- core_vs_edge(xy = grid, boundary = boundary, radius = 2*radius)
grid_edge <- grid[ceo_grid == "edge",,drop = FALSE]

# Use the neighbor-identifying workhorse to identify points in the grid that
# are inside the neighborhoods. Their number is a proxy for neighborhood area.
query_edge <- query[ceo_query == "edge",,drop = FALSE]
d <- fast_distances_up_to_r(query = query_edge,
                           reference = grid_edge,
                           radius = radius)

# Put all the areas together: initialize with full circles,
# then overwrite with the areas calculated for the edge cases:
A <- rep(A0, nrow(query))
A[ceo_query == "edge"] <- pmin(A0,
                              tabulate(d$query_i, nbins = nrow(query_edge)) * step^2)

areas[,i] <- A
print(paste0(i, "/", m))
}

# Output:
list(scales = data.frame(radius = radii, step = steps, error = errors),
     areas = areas)
}

```

4 Example of code behaviour

This uses real data to show how the code works, the times involved, the size of the outputs, etc. In most practical applications, much of this would be wrapped.

```

# This is my local path. You may have to adapt this.
g <- regexpr("fgeo2", getwd())
path <- substr(getwd(), 1, g + attr(g, "match.length"))

# Load data from the Luquillo forest plot, public domain, as an example
load(paste0(path, "data_examples/luquillo.full2.rdata"))
data <- luquillo.full2
dim(data)

```

```
[1] 125446      19
```

```
head(data)
```

	treeID	stemID	tag	StemTag	sp	quadrat	gx	gy	DBHID	CensusID	dbh	pom
1	1	1	-110496	-110496	CECSCH	622	NA	NA	103376		2	NA <NA>
3	2	2	-11392	-11392	PSYBER	503	NA	NA	103377		2	NA <NA>
5	3	3	-114036	-114036	PSYBER	1223	NA	NA	103378		2	NA <NA>
7	4	4	-11756	-11756	CECSCH	901	NA	NA	103379		2	NA <NA>
9	5	5	-12584	-12584	CECSCH	801	NA	NA	103380		2	NA <NA>
11	6	6	-15558	-15558	UNIDEN	303	NA	NA	103381		2	NA <NA>

	hom	ExactDate	DFstatus	codes	nostems	status	date
1	NA	1996-07-03	dead	MAIN;DEADT	1	D	13333
3	NA	1995-02-09	dead	MAIN;DEADT	1	D	12823
5	NA	1996-08-02	dead	MAIN;DEADT	1	D	13363
7	NA	1995-02-23	dead	MAIN;DEADT	1	D	12837
9	NA	1995-02-13	dead	MAIN;DEADT	1	D	12827
11	NA	1995-02-03	dead	MAIN;DEADT	1	D	12817

```
colnames(data)
```

```
[1] "treeID"      "stemID"      "tag"         "StemTag"     "sp"          "quadrat"
[7] "gx"         "gy"         "DBHID"       "CensusID"    "dbh"         "pom"
[13] "hom"        "ExactDate"   "DFstatus"    "codes"       "nostems"     "status"
[19] "date"
```

```
summary(data)
```

treeID	stemID	tag	StemTag
Min. : 1	Min. : 1	Min. : -172721	Min. : -114036
1st Qu.: 31363	1st Qu.: 29528	1st Qu.: 41423	1st Qu.: 28280
Median : 62726	Median : 57530	Median : 84202	Median : 66158
Mean : 62726	Mean : 57731	Mean : 84106	Mean : 64320
3rd Qu.: 94089	3rd Qu.: 86554	3rd Qu.: 124460	3rd Qu.: 97486
Max. : 125450	Max. : 114032	Max. : 175926	Max. : 146191
	NA's : 30956		NA's : 30956

sp	quadrat	gx	gy
Length:125446	Min. : 101	Min. : 0.01	Min. : 0.0
Class :character	1st Qu.: 522	1st Qu.: 86.43	1st Qu.:127.7
Mode :character	Median : 922	Median :164.29	Median :265.2
	Mean : 920	Mean :163.49	Mean :254.9
	3rd Qu.:1316	3rd Qu.:243.75	3rd Qu.:380.8

	Max. :1625	Max. :319.98	Max. :499.9
	NA's :30956	NA's :83870	NA's :83870
DBHID	CensusID	dbh	pom
Min. :103376	Min. :2	Min. : 5.40	Length:125446
1st Qu.:132902	1st Qu.:2	1st Qu.: 15.70	Class :character
Median :160900	Median :2	Median : 27.50	Mode :character
Mean :161103	Mean :2	Mean : 63.61	
3rd Qu.:189924	3rd Qu.:2	3rd Qu.: 79.00	
Max. :217400	Max. :2	Max. :1515.00	
NA's :30956	NA's :30956	NA's :58283	
hom	ExactDate	DFstatus	codes
Min. :0.010	Length:125446	Length:125446	Length:125446
1st Qu.:1.300	Class :character	Class :character	Class :character
Median :1.300	Mode :character	Mode :character	Mode :character
Mean :1.305			
3rd Qu.:1.300			
Max. :3.900			
NA's :58288			
nostems	status	date	
Min. : 1.000	Length:125446	Min. :12717	
1st Qu.: 1.000	Class :array	1st Qu.:12984	
Median : 1.000	Mode :character	Median :13205	
Mean : 1.207		Mean :13152	
3rd Qu.: 1.000		3rd Qu.:13310	
Max. :29.000		Max. :14013	
NA's :30956		NA's :30956	

```

if(FALSE)
{

# Check behaviour
# It is incredibly fast but fills memory: it returns 1-10M rows per second.

n = 5000 # 100-300 k points for hard cases, big ForestGEO plots
radius = 20 # something realistic, rarely > 100
p_focal = 0.05 # proportion of the forest in the focal class

xy <- cbind(x = runif(n, 0, 500), y = runif(n, 0, 500))
focal_xy <- xy[1:round(p_focal*n),]
reference_xy <- xy

```

```

system.time(d <- fast_distances_up_to_r(focal_xy, reference_xy, radius = radius, approx = 0))
dim(d)
head(d)

# Check:
system.time(b <- boundary_around_xy(xy)) # this will take some time for high N
str(b)

plot(b, col = NA, border = "black", xlim = c(0, 500), ylim = c(0, 500))
plot(b, col = NA, border = "black", xlim = c(-15, 50), ylim = c(-15, 50))

# Example with a complex boundary:
library(sf)

boundary <- sf::st_sfc(
  sf::st_polygon(list(rbind(
    c(100, 100),
    c(100, 400),
    c(250, 400),
    c(250, 250),
    c(400, 250),
    c(400, 100),
    c(100, 100)
  )))
)

xy <- data.frame(xy)

loc <- core_vs_edge(
  xy = xy,
  boundary = boundary,
  radius = 20
)

table(loc)

COL <- rep("gray", ncol(xy))
COL[loc == "out"] <- "gray"
COL[loc == "edge"] <- "red"
COL[loc == "core"] <- "black"

```

```

plot(boundary, col = NA, border = "black", xlim = c(0, 500), ylim = c(0, 500))
points(xy$x, xy$y, cex = 0.1, col = COL)

# It can also work by default, surrounding the provided points to guess the boundary:
loc <- core_vs_edge(xy = xy, radius = 20)
table(loc) # probably none out, because of the buffers

COL <- rep("gray", ncol(xy))
COL[loc == "out"] <- "gray"
COL[loc == "edge"] <- "red"
COL[loc == "core"] <- "black"
plot(xy$x, xy$y, cex = 0.1, col = COL, asp = 1)

#plot(grid_for_edge, cex = 0.1, asp = 1)

#
#
#
# *** generalize and make examples with different abundance metrics

if(FALSE)
{
  # Count the number of neighbors and express that as trees / area:
  d <- fast_distances_up_to_r(focal_xy, reference_xy, radius, approx = 0)
  number_of_neighbors <- tabulate(d$query_i, nbins = nrow(focal_xy))
  barplot(table(number_of_neighbors)) # if random locations: Poisson ~ normal distribution

  # Express as densities, edges corrected:
  length(number_of_neighbors) == length(A)
  neighbors_per_area <- number_of_neighbors / A
  hist(neighbors_per_area)
}
}

```


5 Example of a practical application

This is an example on how to put things together. As you can see, these are things that need to be defined outside the functions:

- filters or classes: pay attention to the order of nested loops, removing “dead” or “prior” trees, unknown coordinates, unknown species, etc.
- summarizing metrics: pay attention to units, possible weighting by distance, express in a per-area basis, etc.

```
# This is my local path. You may have to adapt this.
g <- regexpr("fgeo2", getwd())
path <- substr(getwd(), 1, g + attr(g, "match.length"))

# Load data from the Luquillo forest plot, public domain, as an example
load(paste0(path, "data_examples/luquillo.full2.rdata"))
data <- luquillo.full2
dim(data)
```

```
[1] 125446      19
```

```
head(data)
```

	treeID	stemID	tag	StemTag	sp	quadrat	gx	gy	DBHID	CensusID	dbh	pom
1	1	1	-110496	-110496	CECSCH	622	NA	NA	103376		2	NA <NA>
3	2	2	-11392	-11392	PSYBER	503	NA	NA	103377		2	NA <NA>
5	3	3	-114036	-114036	PSYBER	1223	NA	NA	103378		2	NA <NA>
7	4	4	-11756	-11756	CECSCH	901	NA	NA	103379		2	NA <NA>
9	5	5	-12584	-12584	CECSCH	801	NA	NA	103380		2	NA <NA>
11	6	6	-15558	-15558	UNIDEN	303	NA	NA	103381		2	NA <NA>
	hom	ExactDate	DFstatus		codes	nostems	status	date				
1	NA	1996-07-03	dead	MAIN;DEADT		1	D	13333				
3	NA	1995-02-09	dead	MAIN;DEADT		1	D	12823				
5	NA	1996-08-02	dead	MAIN;DEADT		1	D	13363				
7	NA	1995-02-23	dead	MAIN;DEADT		1	D	12837				
9	NA	1995-02-13	dead	MAIN;DEADT		1	D	12827				
11	NA	1995-02-03	dead	MAIN;DEADT		1	D	12817				

```
colnames(data)
```

```

[1] "treeID"      "stemID"      "tag"          "StemTag"      "sp"           "quadrat"
[7] "gx"          "gy"          "DBHID"        "CensusID"     "dbh"          "pom"
[13] "hom"         "ExactDate"   "DFstatus"     "codes"        "nostems"      "status"
[19] "date"

```

```
summary(data)
```

```

      treeID      stemID      tag      StemTag
Min.   :    1  Min.   :    1  Min.   : -172721  Min.   : -114036
1st Qu.: 31363 1st Qu.: 29528 1st Qu.:  41423 1st Qu.:  28280
Median : 62726 Median : 57530 Median :  84202 Median :  66158
Mean   : 62726 Mean   : 57731 Mean   :  84106 Mean   :  64320
3rd Qu.: 94089 3rd Qu.: 86554 3rd Qu.: 124460 3rd Qu.:  97486
Max.   :125450 Max.   :114032 Max.   : 175926 Max.   : 146191
      NA's   :30956      NA's   :30956

      sp      quadrat      gx      gy
Length:125446  Min.   : 101  Min.   :  0.01  Min.   :  0.0
Class :character 1st Qu.: 522 1st Qu.: 86.43 1st Qu.:127.7
Mode  :character Median : 922 Median :164.29 Median :265.2
      Mean   : 920 Mean   :163.49 Mean   :254.9
      3rd Qu.:1316 3rd Qu.:243.75 3rd Qu.:380.8
      Max.   :1625 Max.   :319.98 Max.   :499.9
      NA's   :30956 NA's   :83870 NA's   :83870

      DBHID      CensusID      dbh      pom
Min.   :103376  Min.   : 2  Min.   :  5.40  Length:125446
1st Qu.:132902 1st Qu.: 2 1st Qu.: 15.70  Class :character
Median :160900 Median : 2 Median : 27.50  Mode  :character
Mean   :161103 Mean   : 2 Mean   : 63.61
3rd Qu.:189924 3rd Qu.: 2 3rd Qu.: 79.00
Max.   :217400 Max.   : 2 Max.   :1515.00
NA's   :30956  NA's   :30956 NA's   :58283

      hom      ExactDate      DFstatus      codes
Min.   :0.010  Length:125446  Length:125446  Length:125446
1st Qu.:1.300  Class :character  Class :character  Class :character
Median :1.300  Mode  :character  Mode  :character  Mode  :character
Mean   :1.305
3rd Qu.:1.300
Max.   :3.900
NA's   :58288

      nostems      status      date
Min.   : 1.000  Length:125446  Min.   :12717
1st Qu.: 1.000  Class :array    1st Qu.:12984

```

Median :	1.000	Mode :	character	Median :	13205
Mean :	1.207			Mean :	13152
3rd Qu.:	1.000			3rd Qu.:	13310
Max. :	29.000			Max. :	14013
NA's :	30956			NA's :	30956

```
#
```

6 The ring case

Neighborhoods can be defined as circles (everything from the focal point to a given distance) or rings (things further away than “r” but closer than “R” to the focal point). Rings are conceptually more explicit about the scale at which those measurements apply. If you are doing this type of analysis you should know which option (circular vs ring-shaped neighborhoods) you prefer.

7 Comparing observations with expectations

I present examples on how to compare a given point pattern with the expectations.

First, I present the Ripley’s K and Besag’s L case, using standard R implementation. These are based on counts and Poisson point pattern as the null expectation (that is, independent and uniformly random distribution of points). This functionality was in the original CTFS package but is not implemented directly here as part of the functions in fgeo2.

Second, I present a template of a more general approach based on numerical null models

7.1 Example of Ripley’s K

fgeo2 does not contain Ripley’s K and there is no need for it, as this is a standard spatial analytical tool that is well-maintained in other packages. Instead, I include an example on how you would implement this analysis using {x, y} coordinates.

```
# Example for the living trees of one species
# At what scales is Farama occidentalis in BCI spatially aggregated?

sp = "faraoc"
```

7.2 Example of numerical null model

Numerical null models are very flexible, and could be used for the same purposes as Ripley's K, with added advantages:

- Ripley's K uses the simplest null model (uniform random distribution = Poisson point pattern), but in a customizable numerical null model the user could incorporate different null models, like the torus translation or constrained torus translation.
- The metric of interest in a null model approach could be any, not just counts of neighbors as implicit in Ripley's K. It could be biomass, sum of basal area, etc.
- Null models already incorporate the expectation of fewer neighbors for focal points near the edges, so they can be a direct way of correcting for edge effects for boundaries or arbitrary shape.

In general, null models have nice properties for generalization and flexibility. Their main practical limitations are: (1) the tremendous responsibility on the definition of a reasonable / useful expectation, and (2) computation and storage. If you are afraid of responsibility, do not use null models.

```
# Null models compare observations (one instance) with expectations (many instances).  
# If reality does not deviate much from the expectations by the null model, then  
# the researcher cannot reject the null hypothesis modelled by the null model.  
# If reality deviates from the expectations, then we can reject the null hypothesis.
```