

Evaluation of a set of functions related to Ripley's K analysis

Gabriel Arellano

2026-05-27

Contents

1	Goal	1
2	Overall criteria	2
3	Preliminaries for code evaluation	2
4	Review of original code	4
5	Design principles for alternative code	5
6	Current status of alternative code	6

1 Goal

Evaluate, understand, and decide over a set of related functions that count individuals in circular areas and compare those counts with a null expectation, implementing Ripley's K test. At the moment, this is an isolated set of functions; see "1-dependencies-and-priorities-2026-05-20.pdf" report.

This set contains the following functions:

- `CalcRingArea` in `spatial > RipUvK.r`
- `partialcirclearea` in `spatial > RipUvK.r`
- `circlearea` in `spatial > RipUvK.r`
- `Annuli` in `spatial > RipUvK.r`
- `NDcount` in `spatial > NeighborDensityFun.r`

- `NeighborDensities` in `spatial` > `NeighborDensityFun.r`
- `RipUvK` in `spatial` > `RipUvK.r`

2 Overall criteria

These functions are located in two different scripts but all are related to the same type of calculation: counting number of neighbors in a given area, and compare that with some null expectation (Ripley's K function). The set covers all code in `RipUvK.r` and `NeighborDensityFun.r` original scripts.

This type of calculation is widely used in ForestGEO / ATFS, for example to estimate competition around target individuals and so on. Perhaps not to much for point pattern analyses. But, in general, counting points in certain areas around target individuals (or locations) is valuable. Perhaps most of these functions can be incorporated into a more general function that counts (or summarizes) data belonging to different classes, within areas that have a minimum and a maximum radius.

Modelling or comparison with null expectations is less valuable for the `fgeo2` package. The Ripley's K function is surely maintained by other packages: `spatstat::Kest` and `spatstat::Lest` may be the gold standard (and fastest?). The code implements a well-known mathematical formula so it should not require a lot of maintenance. However, the `RipUvK` function is a clear candidate to remove, and substitute by a more general tutorial on how to apply Ripley's K to forest data.

The study of neighborhoods can be done in many different ways, so it would be important to maximize flexibility in things like selecting the focal trees, picking neighborhood-level metrics, etc. If such flexibility is not feasible to incorporate within `fgeo2` code, then it should be left out the functions, and substitute by examples or guidelines that the user can apply to their case and goals.

3 Preliminaries for code evaluation

First, we declare the path to `fgeo2` folder and all the functionality needed:

```
# This is my local path. You may have to adapt this.
g <- regexpr("fgeo2", getwd())
path <- substr(getwd(), 1, g + attr(g, "match.length"))

# Source useful functionality:
source(paste0(path, "evaluation/0-code-to-source.r"))
```

```
# Other functionality:  
library(igraph)
```

Warning: package 'igraph' was built under R version 4.5.2

Attaching package: 'igraph'

The following objects are masked from 'package:stats':

decompose, spectrum

The following object is masked from 'package:base':

union

```
library(fgeo)
```

```
-- Attaching packages ----- fgeo 1.1.4 --
```

```
v fgeo.analyze 1.1.15      v fgeo.tool    1.2.10  
v fgeo.plot    1.1.11      v fgeo.x       1.1.4
```

```
-- Conflicts ----- fgeo_conflicts() --  
x fgeo.tool::filter() masks stats::filter()
```

```
# This chunk of code will source all the old R scripts:  
f1 <- list.files(paste0(path, "01_utilities/R/old/"),  
                full.names = TRUE, pattern = "\\..R$", ignore.case = TRUE)  
f2 <- list.files(paste0(path, "02_biomass/R/old/"),  
                full.names = TRUE, pattern = "\\..R$", ignore.case = TRUE)  
f3 <- list.files(paste0(path, "03_spatial/R/old/"),  
                full.names = TRUE, pattern = "\\..R$", ignore.case = TRUE)  
f4 <- list.files(paste0(path, "04_abundance_diversity/R/old/"),  
                full.names = TRUE, pattern = "\\..R$", ignore.case = TRUE)  
f5 <- list.files(paste0(path, "05_demography_growth/R/old/"),  
                full.names = TRUE, pattern = "\\..R$", ignore.case = TRUE)  
f6 <- list.files(paste0(path, "06_topography_habitat/R/old/"),  
                full.names = TRUE, pattern = "\\..R$", ignore.case = TRUE)
```

```
f7 <- list.files(paste0(path, "07_mapping_plotting/R/old/"),
                full.names = TRUE, pattern = "\\R$", ignore.case = TRUE)

ff <- c(f1, f2, f3, f4, f5, f6, f7)
for(f in ff) source(f)
```

4 Review of original code

```
# Use this to see the source code, manually
if(FALSE)
{
  CalcRingArea
  partialcirclearea
  circlearea
  Annuli
  NDcount
  NeighborDensities
  RipUvK
}
```

Some functions work on areas, other functions work on counts.

The function `circlearea` is trivial.

The `partialcirclearea` makes edge corrections considering the dimensions of the plot. That will be important in forest applications. It seems the most interesting piece of functionality. Apparently Condit implemented Goreaud & Pélissier (1999, Journal of Vegetation Science 10: 433-438), which are analytical solutions for circular and rectangular study areas. I consider keeping it, although it is very cryptic and virtually impossible to “touch”.

`CalcRingArea` seems to be just the accumulated sum of (possibly truncated) circles, one for each point in a dataset. Note:

- it calculates areas of circles, not rings, despite its name
- it does not correct for overlap, so it counts many times the same areas

The area of rings (“annuli”) is calculated by the `Annuli` function, which repeatedly calls `CalcRingArea` and subtracts small circles from larger circles successively. This general approach should be kept.

`NDcount` and `NeighborDensities` try to do similar things. `NDcount` (author Tania Brenes) wraps `splancs::khat` function to count neighbors with edge correction. `NeighborDensities`

(author Rick Condit) incorporates filters, the possibility of counting neighbors from a different dataset (e.g. adult trees neighboring seedlings), and “counting” as accumulated basal area instead of individual counts. It depends on specific column names like “status”. There are risky unit conversions inside, it does make strong assumptions about the units of the input: clearly a vulnerable point.

`RipUvK` is the highest-level function in this group of functions. It runs over each species calculating a series of spatial statistics for con-specifics at different distances. It seems easy to use, it is not very long. It is not very flexible, and calls external functions of undefined packages, probably `splancs` package: `khat`, `spoints`.

Because the code has obvious points of vulnerability (units, column names, ...), and has very limited flexibility (filters, aggregation, expectations), I decide to rewrite.

5 Design principles for alternative code

Ideally, replacement code should incorporate the flexibility to count neighbors of any class to any class, for example: adult tree neighboring seedlings, individuals of species A neighboring species A (conspecifics), individuals of species A neighboring species B (some heterospecifics), etc. The conspecific vs. heterospecific counting, which is the most frequent, is just a special case. But parameters should be general, not specific assuming this particular application.

It should be general and count everything in the input. If the user wants to focus on trees that are alive, for example, it should subset the dataset first. It does not seem robust to filter from within the function and make assumptions about column names, etc. `Filter outside` gives all the flexibility to the user, even if they keep lianas and trees together, or whatever they want. This is aligned with the goal of reducing dependency from “orthodox CTFS format”.

Edge effects are important. Original code handles rectangular boundaries only. Ideally, we should be able to define edges of any arbitrary shape. For example, so Ngel Nyaki plot can truncate circles to the border of the forested part of the plot, not just the rectangular borders. Or do counts + areas for specific habitat types within the rectangular plot, etc. Goreaud & Pélissier (1999) proposed an algorithm to generalize their approach to arbitrary boundaries, but implementing that would involve “fastidious calculations” (as they say) and would be prone to error. A grid-based approach approaches brute force but is easier to code and check. It should be preferred if computational costs allow it.

Calculating distances to focal points could be relatively easy, but memory and computation is important for large datasets.

It should be unit agnostic: return everything on the same units as the input.

Ideally, the code should be fast. Solutions without loops would be preferred. Memory could be a constrain so tutorials should teach how to do things by chunks.

Ideally, we want to separate the identification of neighbors (who is neighbor of whom) from the counting. Counting is just a metric of abundance, but we could use accumulated basal area or biomass, etc. Both steps are somehow mixed in the current code, limiting flexibility.

Comparing with expectations: the current code uses Ripley's K function, which is a statistic that measures the deviation of a given count in a given area relative to the null expectation of random placement (Poisson point process or PPP). This is a well-known technique implemented in many other packages. It should be deprecated and substituted by tutorials on how to apply it. The tutorials can also show how to implement numerical null models to compare observed counts (or abundances) of neighbors around focal trees with the expectations. Numerical null models would have clear advantages:

- They would not be limited to individual counts as now, but could incorporate observations vs. expectations for other metrics, such as accumulated basal area, total biomass, etc.
- The user could model, potentially, other null expectations, not just PPP.
- If the null model is constrained to a given boundary, it would reflect edges, so deviations from the model would solve for edge effects directly.

In any case, numerical null models require fast code, because of computational costs. Because these are very flexible tools, they should not be hard-coded in fgeo functionality.

If all these things are implemented, an updated workflow would clearly expand the current feasible workflow:

- allow the use of any metric of abundance
- neighbors of any class around focal points of any class, defined outside
- filters are outside: maximum flexibility (and responsibility)
- minimize assumptions: work on $\{x, y\}$ general datasets, with boundaries expressed as spatial objects
- account for edge effects in boundaries of arbitrary shape
- allow the use of any null distribution of points, not just PPP

6 Current status of alternative code

See “neighborhoods.qmd” in the spatial module.